

# Airflow in Practice:

## How I Learned to Stop Worrying and Love DAGs



Sarah Schattschneider  
PyBay 2019

# Who Am I?

- Software Engineer @ Blue Apron
  - Blue Apron - meal kit delivery service to make incredible home cooking accessible to everyone
  - Data Operations Team
- Airflow Interests:
  - Extract Transform Load (ETL) workflows
  - BigQuery Permissioning
  - Testing
- @sschatts



# Outline

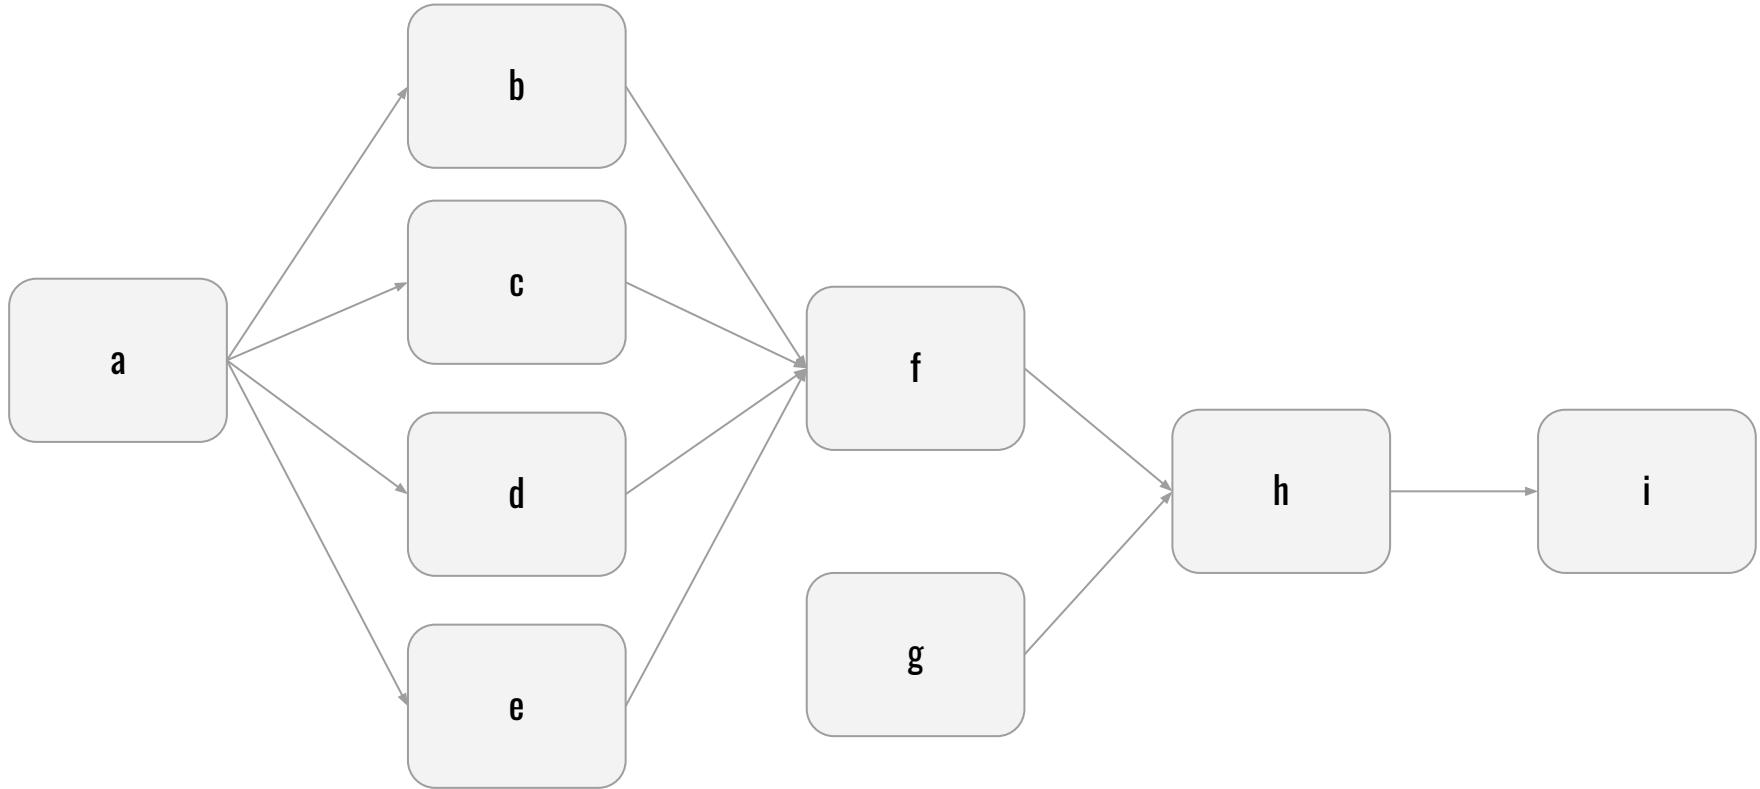
- Background
- Airflow Concepts
- Evaluating Airflow
- Challenges of Airflow
- Common Use Cases
- Best Practices

**Background**

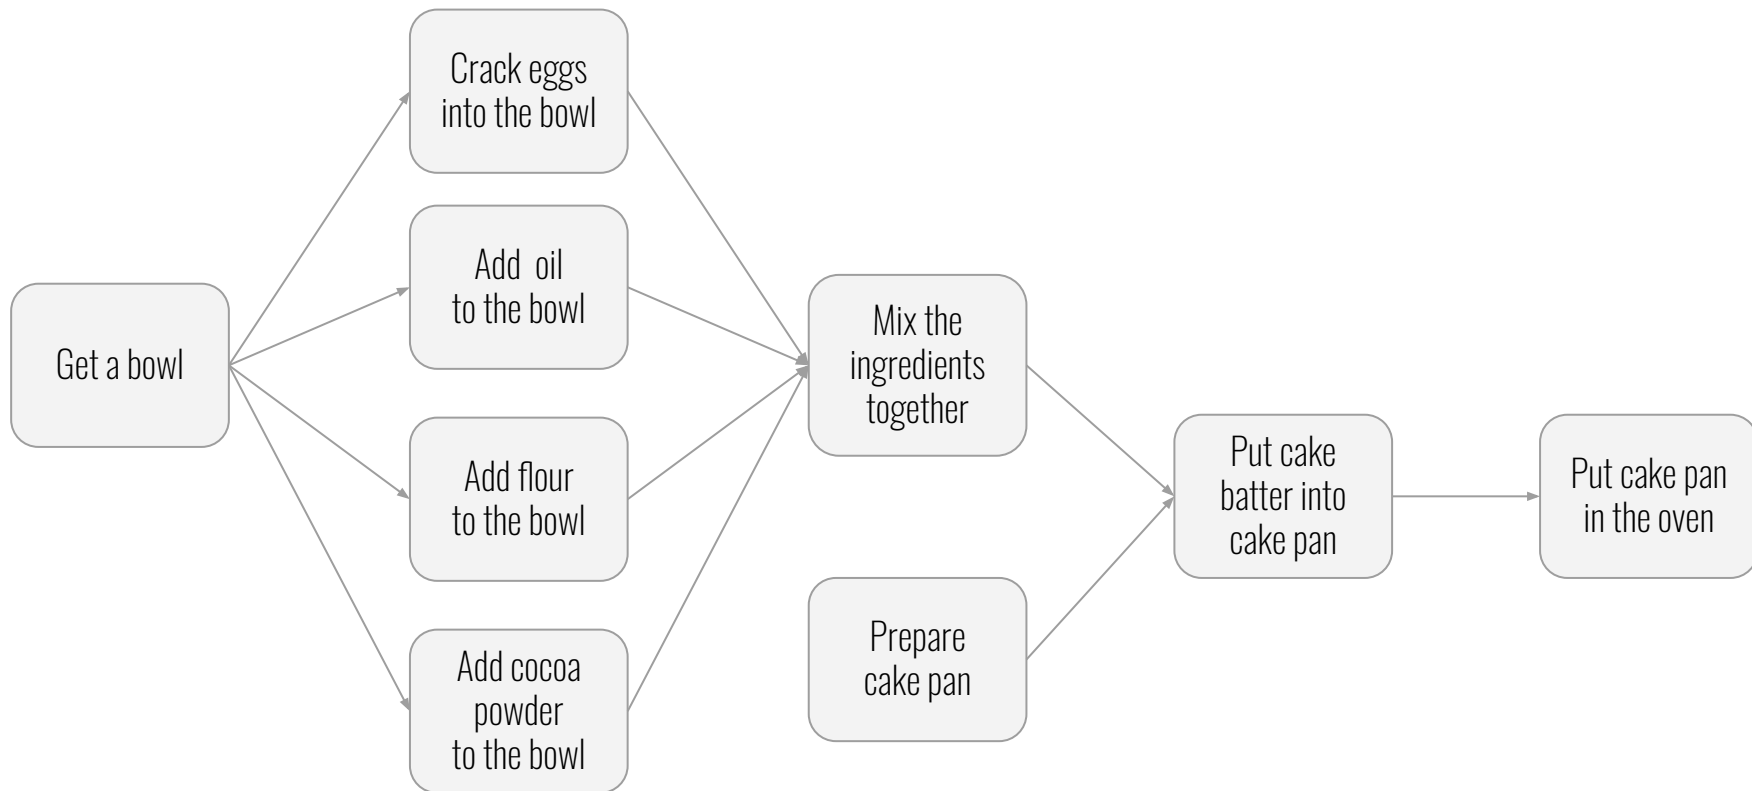
# Once upon a time there was cron

```
0 1 * * * /bin/bash echo "hello world"  
0 3,15 * * * /scripts/script.sh  
0 2 * * mon /scripts/weekly.sh
```

# Directed Acyclic Graph (DAG)



# Directed Acyclic Graph (DAG)



# Airflow Concepts



# What is Airflow?

- Programmatically author workflows
- Stateful scheduling
- Rich Command Line Interface (CLI) and User Interface (UI)
- Logging, monitoring, and alerting
- Modularity and Testability
- Solves common problems in batch processing

# Declare a DAG

```
def hello_world():
    return 'Hello World!'

default_args = {
    'owner': 'airflow',
    'depends_on_past': False,
    'start_date': datetime(2019, 8, 1),
    'email_on_failure': False,
    'retries': 2,
    'max_active_runs': 1
}


dag = DAG('example_dag', default_args=default_args, schedule_interval='0 21 * * *')
dag.doc_md = """
##### DAG Summary
This is an example dag for my PyBay presentation

##### Points of Contact
Data Engineering
Slack channel: #data-engineering
"""


start = DummyOperator(task_id='start', dag=dag)
start.doc_md = """ #start operator """

hello = PythonOperator(
    task_id='hello',
    python_callable=hello_world,
    dag=dag
)
hello.doc_md = """ #hello world operator """
start >> hello
```

# User Interface

 Airflow

DAGs Data Profiling ▾ Browse ▾ Admin ▾ Docs ▾ About ▾ SQL ▾

21:47 UTC 

On **DAG: example\_dag** schedule: 0 21 \* \* \* Graph View  Tree View  Task Duration  Task Times  Landing Times  Gantt  Details  Code  Refresh

### DAG Summary


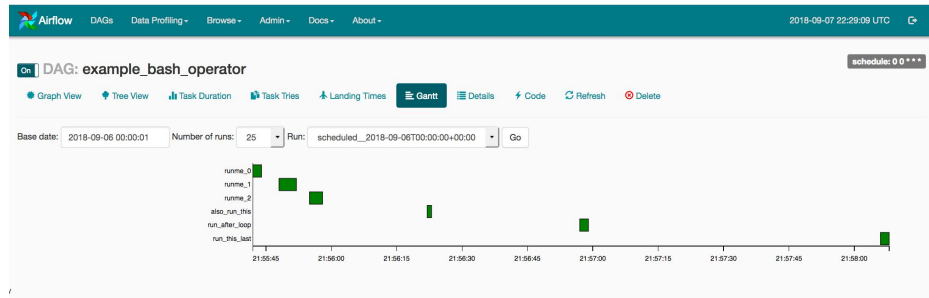
This is an example dag for my PyBay presentation

### Points of Contact

Data Engineering Slack channel: #data-engineering

success Run:  Layout:   DummyOperator PythonOperator success running failed skipped retry queued no status 

# Metrics



Airflow README

# Hooks

Encapsulate all  
interactions with an API

# Operators

Instantiate a hook or two to  
complete a single goal

# Tasks

An instance of an operator

# Hooks

Cake Pan



Oven



# Operators

Put Cake Pan in Oven



# Tasks

```
CakePanToOvenOperator(  
  task_id='cake_to_oven',  
  cake_ingredients={  
    'eggs',  
    'oil',  
    'cocoa powder',  
    'flour'  
  },  
  oven_conn_id=OVEN_CONN_ID,  
  cake_pan_conn_id=CAKE_PAN_CONN_ID,  
  dag=dag  
)
```

Put cake pan  
in the oven

# Hooks



Google Cloud Storage



Google BigQuery

# Operators



Google Cloud Storage



Google BigQuery

# Tasks

```
gcs_to_bq =
GoogleCloudStorageToBigQueryOperator(
    task_id='gcs_to_bq',
    bucket='my_bucket',
    source_objects=['path/to/my_file.csv'],
    destination_project_dataset_table='dataset.table',
    schema=transfer_schema,
    write_disposition='WRITE_TRUNCATE',
    skip_leading_rows=1,
    google_cloud_storage_conn_id=GCS_CONN_ID,
    bigquery_conn_id=BQ_CONN_ID,
    dag=dag
)
```

# Airflow Hooks

- AwsHook
- JdbcHook
- GoogleCloudStorageHook
- CassandraHook
- BigQueryHook
- RedisHook
- DatadogHook
- AzureFileShareHook
- GrpcHook
- JiraHook
- SSHHook
- HiveCliHook
- SlackHook
- DockerHook
- S3Hook
- MySqlHook
- PrestoHook

and many more...



# Airflow Operators

- PythonOperator
  - BashOperator
  - EmailOperator
  - S3ToRedshiftTransfer
  - RedshiftToS3Transfer
  - BigQueryToCloudStorageOperator
  - GoogleCloudStorageToBigQueryOperator
  - SlackAPIPostOperator
  - MySqlToHiveTransfer
  - SimpleHttpOperator
  - DataprocOperationBaseOperator
  - JiraOperator
  - DockerOperator
  - PostgresOperator
- and many more...

# Evaluating Airflow

# What value does Airflow add?

- Retries tasks elegantly
- Alerts on failure
- Re-run specific tasks in the DAG
- Supports distributed execution

# Third Party Data Integrations



Google Cloud Platform



and many more...

# Cron Scheduling

**Monday at 4:00**

0

minute

4

hour

\*

day  
(month)

\*

month

1

day  
(week)

---

**10th day of the month  
at 10:05**

5

minute

10

hour

3

day  
(month)

\*

month

\*

day  
(week)

# Cron Scheduling

- Time zone safety
- Flexibility of determining when to run

**Does Airflow Have an Ugly Side?  
How to Overcome these Challenges?**

# Upgrades

- Upgrades can be more challenging when you have custom hooks and operators





## Help with Upgrade Challenges:

- Having DAGs run in staging environment
- Better unit tests
- Contribute to upstream regularly





# Data Sharing

- Xcoms - Airflow cross-communication between tasks

 Task Instance Details	 Rendered Template	 Log	 XCom
XCom			
Key		Value	
return_value		b'\x80\x03x\x0c\x00\x00\x00Hello World!q\x00.'	

- Airflow Variable

List (1)	Create	Add Filter▼	With selected▼	Search: key, val
<input type="checkbox"/>		Key	Val	
<input type="checkbox"/>	 	foo	'hello world'	

- Python Variable
- Environment Variable

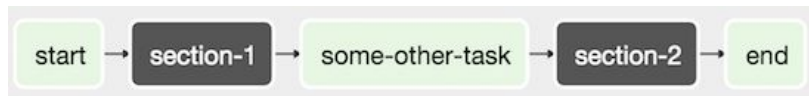
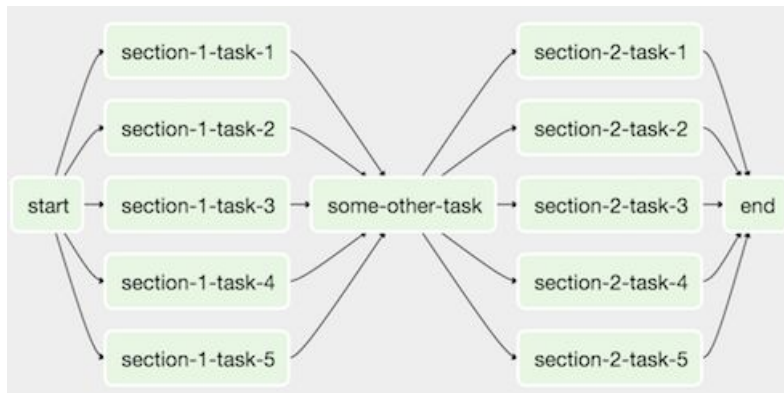
# Common and Fun Use Cases

# Extract Transform Load (ETL) Jobs


- Airflow enables moving data and transforming data very easily
- Can create custom Hooks and Operators for Third Party APIs

# Subdags? How do you use them?

- Used when there are repeating patterns



# Subdags Example

 **Airflow**

DAGs

Data Profiling ▾

Browse ▾

Admin ▾

Docs ▾


About ▾


SQL ▾


18:10 UTC


⏻


On **DAG: snapshot-pss** schedule: 0 \*\*\*\*


 **Graph View**


 Tree View


 Task Duration


 Task Tries

 Landing Times

 Gantt

 Details


 Code

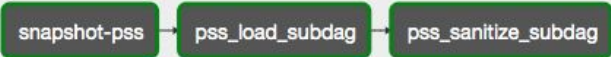
 Refresh

success Run:  Layout:

**SubDagOperator**

success running failed skipped retry queued no status



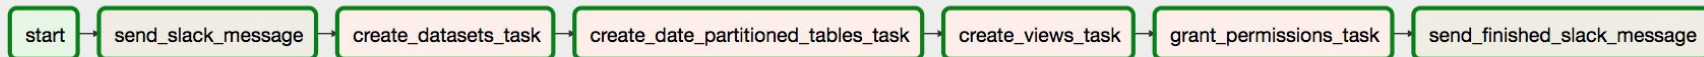


# Slack Integrations

- Notify different Slack channels

```
slack_notification= SlackAPIPostOperator(  
    task_id='slack_notification',  
    username='Airflow Alerting User',  
    slack_conn_id=SLACK_CONN_ID,  
    channel='#airflow-alert-channel',  
    text='Hello World',  
    dag=dag  
)
```

# BigQuery Permissioning



# ML Feature Extraction

- Export data from our data warehouse
- Use the export data in Spark to train model



# Best Practices

# Testing

- Unit tests
  - Supporting Python functions
  - Custom hooks and operators
- Acceptance test to run ``airflow list_dags``

# Doc MD for the DAG

## #### *DAG Summary*

SHORT 2-3 sentence paragraph of project description goes here.

## #### *Mission Critical*

Yes/No and a short blurb for context

## #### *On Failure Actions*

Rerun DAG? Yes/No and a short blurb for context

## #### *Points of Contact*

Multiple points of contact are preferable with area of responsibility and title if possible. If updating, make sure to include the date

## #### *Business Use Case/Process*

<https://www.lucidchart.com/documents/edit/12345>

## #### *Prerequisites/Dependencies/Resourcing*

- Some cool documentation

# In the Codebase

```
dag = DAG('example_dag', default_args=default_args, schedule_interval='0 21 * * *')
```

```
dag.doc_md = """
```

```
#### DAG Summary
```

```
SHORT 2-3 sentence paragraph of project description goes here.
```

```
#### Mission Critical
```

```
Yes/No and a short blurb for context
```

```
#### On Failure Actions
```

```
Rerun DAG? Yes/No and a short blurb for context
```

```
#### Points of Contact
```

```
Multiple points of contact are preferable with area of responsibility and title if possible. If updating, make sure to include the date
```

```
#### Business Use Case/Process
```

```
https://www.lucidchart.com/documents/edit/12345
```

```
#### Prerequisites/Dependencies/Resourcing
```

- Some cool documentation

```
"""
```

# Materialized in the UI

The screenshot displays the Apache Airflow web interface. At the top, a teal navigation bar contains the Airflow logo and links for DAGs, Data Profiling, Browse, Admin, Docs, About, and SQL. The current time is 21:25 UTC. Below the navigation bar, the main header shows 'On DAG: example\_dag' with a 'schedule: 0 21 \* \* \*' indicator. A row of view options includes Graph View (selected), Tree View, Task Duration, Task Tries, Landing Times, Gantt, Details, Code, and Refresh. The DAG Summary section contains the following text: 'This is an example dag for my PyBay presentation', 'Mission Critical' (with a note 'No this is not mission critical'), 'On Failure Actions' (with a note 'The latest run of this DAG will catch up, no need to re-run'), and 'Points of Contact' (with a note 'Data Engineering Slack channel: #data-engineering'). Below the summary, there are controls for 'Run' (a dropdown), 'Layout' (set to 'Left->Right'), and a 'Go' button. A search bar is also present. The task list at the bottom shows 'BashOperator' and 'DummyOperator' with status indicators: success (green), running (green), failed (red), skipped (pink), retry (yellow), queued (blue), and no status (grey). The DAG graph at the bottom shows a simple flow from 'start' to 'hello'.

Airflow DAGs Data Profiling Browse Admin Docs About SQL 21:25 UTC

On DAG: example\_dag schedule: 0 21 \* \* \*

Graph View Tree View Task Duration Task Tries Landing Times Gantt Details Code Refresh

**DAG Summary**  
This is an example dag for my PyBay presentation  
**Mission Critical**  
No this is not mission critical  
**On Failure Actions**  
The latest run of this DAG will catch up, no need to re-run  
**Points of Contact**  
Data Engineering Slack channel: #data-engineering

Run: Layout: Left->Right Go Search for...

BashOperator DummyOperator success running failed skipped retry queued no status

start → hello

# Thanks!

...

@sschatts  
sarah@blueapron.com



Software Engineer - Data Operations  
Blue Apron